Database Management Systems

BITS Pilani
Hyderabad Campus

Dr.R.Gururaj
CS&IS Dept.

Lecture Session-11

Data Storage

innovate achieve lead

Content

- *Disk pack features*
- *Records and Files*
- *File operations*
- *Ordered and unordered features*

Disk Storage

- Disk is the preferred secondary storage device for high storage capacity and low cost.
- Data stored as magnetized areas on magnetic disk surfaces.
- A **disk pack** contains several magnetic disks connected to a rotating spindle.
- Disks are divided into concentric circular **tracks** on each disk **surface**.
  - Track capacities vary typically from 4 to 50 Kbytes or more

Figure 13.1
(a) A single-sided disk with read/write hardware. (b) A disk pack with read/write hardware.

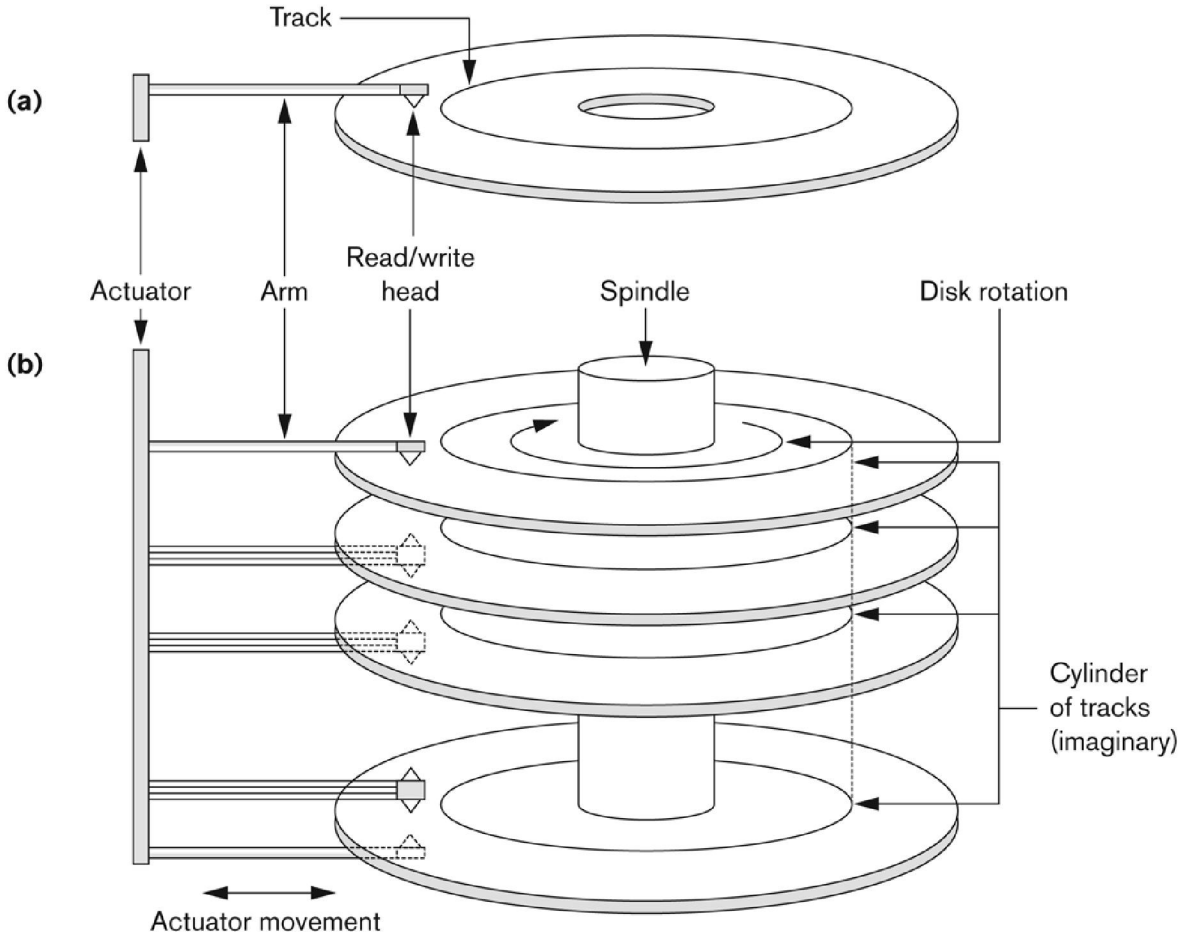

- A track is divided into smaller **blocks** or **sectors.**

- The division of a track into **sectors** is hard-coded on the disk surface and cannot be changed.

- A track is divided into **blocks**.

  1. The block size B is fixed for each system.

     Typical block sizes range from B=512 bytes to B=4096 bytes.

  2. Whole blocks are transferred between disk and main memory for processing.

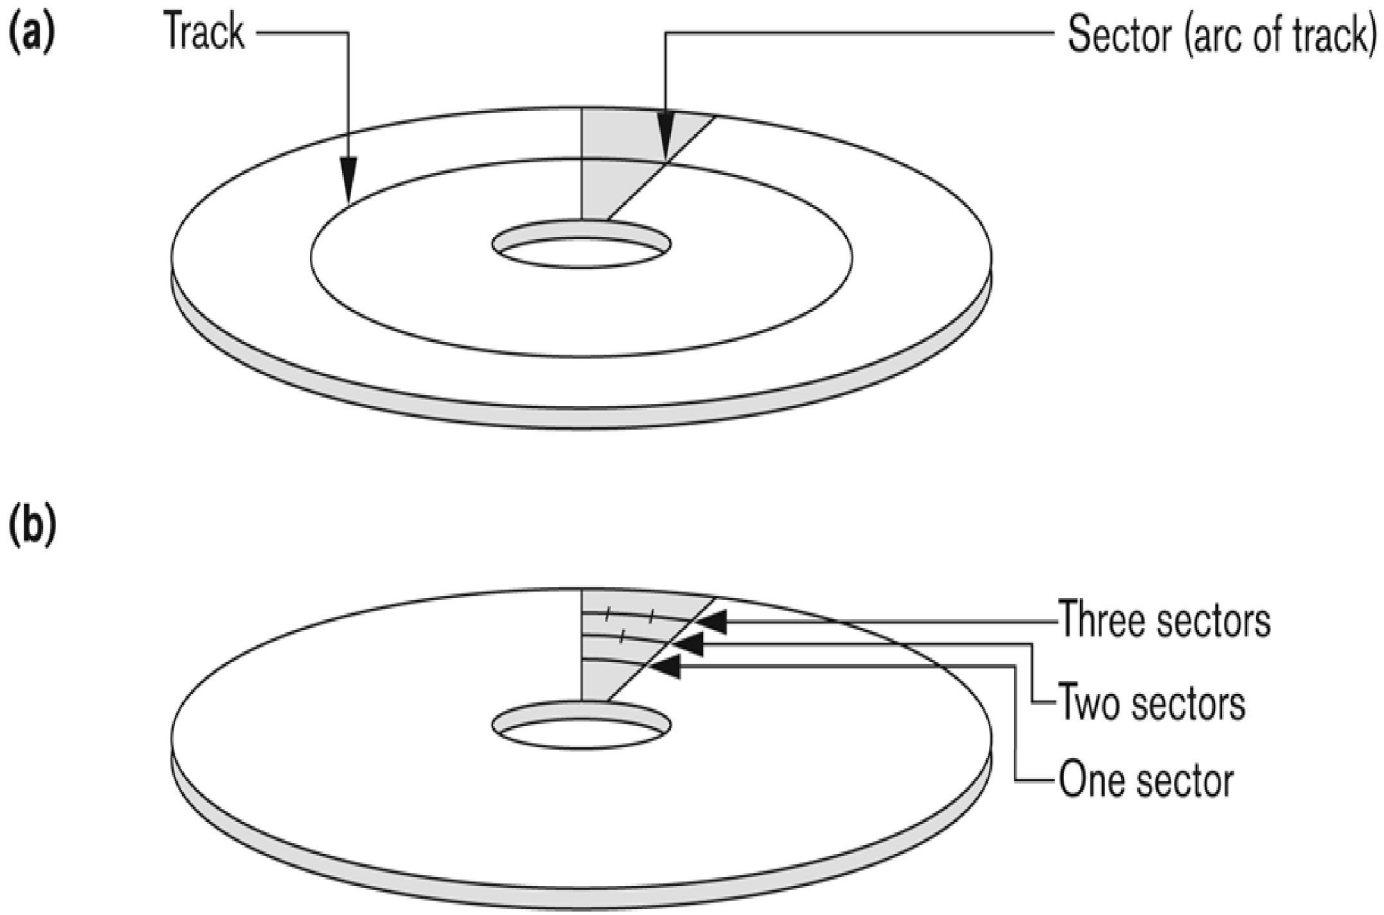

Figure 13.2 Different sector organizations on disk. (a) Sectors subtending a fixed angle. (b) Sectors maintaining a uniform recording density.

- A **read-write head** moves to the track that contains the block to be transferred.

  Disk rotation moves the block under the read-write head for reading or writing.
- A physical disk block (hardware) address consists of:
  - a cylinder number (imaginary collection of tracks of same radius from all recorded surfaces)
  - the track number or surface number (within the cylinder)
  - and block number (within track).
- Reading or writing a disk block is time consuming because of the seek time s (time to position the head on required track)

  3-7msec and rotational delay (latency) – time to position at the beginning of the required block **rd**.

  3-4 msec with 15000rpm

Block transfer time. Smaller than above two.

Files and Records

- A **file** is a *sequence* of records, where each record is a collection of data values (or data items).
- A **file descriptor** (or **file header**) includes information that describes the file, such as the *field names* and their *data types*, and the addresses of the file blocks on disk.
- Records are stored on disk blocks.
- The **blocking factor** (**bfr)** for a file is the (average) number of file records stored in a disk block.
- A file can have **fixed-length** records or **variable-length** records.

- File records can be **unspanned** or **spanned**
  - **Unspanned**: no record can span two blocks
  - **Spanned**: a record can be stored in more than one block
- The physical disk blocks that are allocated to hold the records of a file can be *contiguous, linked*.
- In a file of fixed-length records, all records have the same format. Usually, unspanned blocking is used with such files.
- Files of variable-length records require additional information to be stored in each record, such as **separator characters** and **field types**.
  - Usually spanned blocking is used with such files.

File operations

Typical file operations include:

- **OPEN**: Readies the file for access, and associates a pointer that will refer to a *current* file record at each point in time.
- **FIND**: Searches for the first file record that satisfies a certain condition, and makes it the current file record.
- **FINDNEXT**: Searches for the next file record (from the current record) that satisfies a certain condition, and makes it the current file record.
- **READ**: Reads the current file record into a program variable.
- **INSERT**: Inserts a new record into the file & makes it the current file record.
- **DELETE**: Removes the current file record from the file, usually by marking the record to indicate that it is no longer valid.
- **MODIFY**: Changes the values of some fields of the current file record.
- **CLOSE**: Terminates access to the file.
- **REORGANIZE**: Reorganizes the file records.

  For example, the records marked deleted are physically removed from the file or a new organization of the file records is created.
- **READ_ORDERED**: Read the file blocks in order of a specific field of the file.

Unordered Files

Also called a *heap* or a *pile* file.

New records are inserted at the end of the file.

A *linear search* through the file records is necessary to search for a record.

- This requires reading and searching half the file blocks on the average, and is hence quite expensive.

Record insertion is quite efficient.

Reading the records in order of a particular field requires sorting the file records.

Ordered Files

- Also called a *sequential* file.
- File records are kept sorted by the values of an *ordering field*.
- Insertion is expensive: records must be inserted in the correct order.

A *binary search* can be used to search for a record on its *ordering field* value.

  - This requires reading and searching $\log_2$ of the file blocks on the average, an improvement over linear search.
- Reading the records in order of the ordering field is quite efficient.

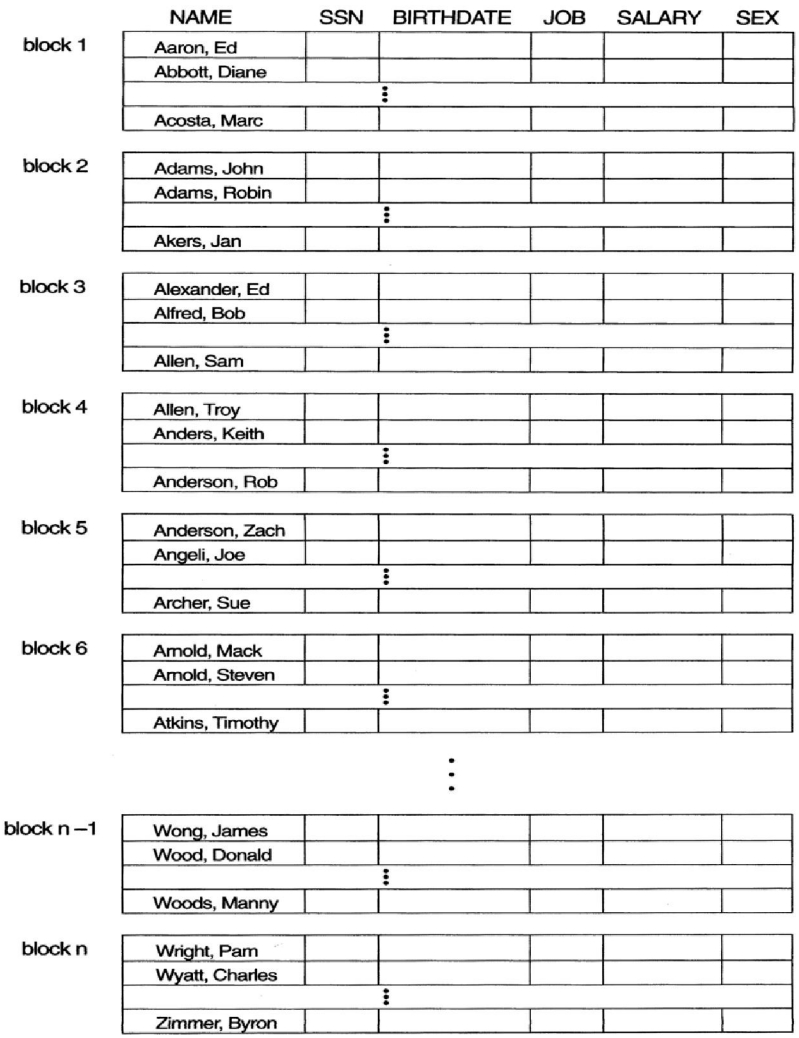
NAME
SSN
BIRTHDATE
JOB
SALARY
SEX
block 1
Aaron, Ed
Abbott, Diane
Acosta, Marc
block 2
Adams, John
Adams, Robin
Akers, Jan
block 3
Alexander, Ed
Alfred, Bob
Allen, Sam
block 4
Allen, Troy
Anders, Keith
Anderson, Rob
block 5
Anderson, Zach
Angeli, Joe
Archer, Sue
block 6
Arnold, Mack
Arnold, Steven
Atkins, Timothy
block n –1
Wong, James
Wood, Donald
Woods, Manny
block n
Wright, Pam
Wyatt, Charles
Zimmer, Byron

Summary

- ✓ *What is Disk storage*
- ✓ *Disk characteristics*
- ✓ *Disk pack structure*
- ✓ *Files and Records*
- ✓ *Ordered and unordered files*