



BITS Pilani
Hyderabad Campus

Database Management Systems

Dr.R.Gururaj
CS&IS Dept.

Lecture Session-17

Concurrency Control



Contents

- ☐ *Introduction to Concurrency Control*
- ☐ *Implementing Serializability*
- ☐ *Lock-based protocols*
- ☐ *Deadlock condition*
- ☐ *Two-phase locking protocol*
- ☐ *Time-phase locking protocol*

Introduction



- ❖ In a DBMS multiple transactions are executed concurrently.
- ❖ If the transactions are executed concurrently then the resources can be utilized more efficiently hence more throughput is achieved.
- ❖ Here, for transactions we consider data items as resources because transactions process data by accessing them.
- ❖ When multiple transactions access data elements in a concurrent way, this may destroy the consistency of the database.

Implementing Serializability



One way to ensure *serializability* is to allow the transactions to access the data items in a mutually exclusive manner.

This is to make sure that when one transaction access a data item no other transaction can modify that data item.

The following techniques implement mutual exclusion and control concurrency.

1. *Lock-based protocols*
2. *Timestamp-based protocols*



1. Concurrency Control Using Locks: A data item may be locked in various modes.

i) **Shared** (denoted by *S*): if a transaction obtains a shared mode lock on a data item *Q*, it can read *Q* but not modify *Q*.

(ii) **Exclusive** (denoted by *X*): if this lock is obtained, a transaction can read or write the data item.



Lock Compatibility Matrix

	S	X
S	True	False
X	False	False

This says that if a transaction T_i obtains a lock on a data item in *S-mode*, other transaction can get a lock on the same item in *S-mode* but not in *X-mode*.

If a transaction obtains a lock in *X-mode* on a data item no other transaction can obtain a lock on the same data item in any mode.

Deadlock



The Mutual exclusion mechanism leads to deadlock situation.

For example, if transaction T_i holds a lock on a data item (Q) in *X-mode* and waits for a lock on another data item (P) which is locked by another transaction T_j in *X-mode*, further to release the lock on P, T_j must acquire a lock on Q, which is locked by T_i . This is a circular wait condition and results in a *deadlock* situation.



Wait-for Graph

Deadlock condition can be determined by a *wait-for* graph.

All transactions of the schedule become vertices.

And we have an edge between two transactions T_i and T_j if T_i is waiting for T_j to release a lock on a data item.

If the graph has a cycle then we can say that the schedule will result in a deadlock.


T ₁	T ₂	T ₃	T ₄
S(A)			
R(A)			
	X(B)		
	W(B)		
S(B)			
		S(C)	
		R(C)	
	X(C)		
			X(B)
		X(A)	

S(A) means transaction locks A
in share mode

R(A) – transaction reads A

X(C) – Transaction locks
C in X-mode

W(B) – transaction write B





In the above graph there exists a cycle hence this schedule leads to deadlock.

If a transaction T_i requests a lock and transaction T_j holds a conflicting lock. The lock manager can use one of the following *policies to prevent deadlocks*.

Timestamp based:

Wait-Die: If T_i is older than T_j it is allowed to wait otherwise aborted.

Wound-wait: If T_i older than T_j allowed to run by aborting T_j else T_i will wait.

Priority based

Wait-Die: If T_i has higher priority than T_j it is allowed to wait otherwise aborted.

Wound-wait: If T_i lies higher priority it is allowed to run by aborting T_j else T_i will wait.



Two-phase locking protocol:

This protocol answers serializability.

According to this each transaction issues lock and unlock requests in two phases.

- i) *Growing phase*: In this phase, a transaction may obtain locks but may not release any lock.
- ii) *Shrinking phase*: In this phase, a transaction may release locks but may not obtain any new locks.

The two-phase locking protocol ensures conflict serializability.

It does not ensure freedom from deadlock.



2. Timestamp-based Concurrency Control

Maintaining the ordering between every pair of conflicting transactions is significant.

If we select the ordering in advance, we can achieve serializability. Time-stamping is a method to fix the ordering.

Each transaction is assigned a unique fixed timestamp.

If $TS(T_i) < TS(T_j)$, this implies that T_i should be executed before T_j .



The time-stamps determine the serializability order.
Each data item is associated with two timestamp values.

W-timestamp(Q) – represents the largest timestamp of any transaction that successfully executes Write(Q).

R-timestamp(Q) - which denotes the largest time stamp of any transaction that successfully executed Read(Q).

These values are updated whenever read(Q) or write(Q) are executed.



Timestamp ordering Protocol:

This protocol operates as follows:

i) Suppose Transaction T_i issues read(Q)

If $TS(T_i) < W\text{-stamp}(Q)$, then it implies that T_i need to read Q which was already overwritten.

Hence read operation is rejected and T_i is rolled back.

If $TS(T_i) \geq W\text{-timestamp}(Q)$ then read operation is executed.

ii) Suppose T_i issues write (Q)

If $TS(T_i) < R\text{-timestamp}(Q)$ it implies that the value of Q being produced by T_i had to be written long back.

Hence reject T_i & roll back.

If $TS(T_i) < W\text{-timestamp}(Q)$, T_i is attempting to write some absolute value of Q.

Hence reject T_i & roll back.

Otherwise write operation is executed.



Summary

- ✓ *Concepts related to Concurrency Control*
- ✓ *Approaches for Implementing Serializability*
- ✓ *How lock-based protocols work*
- ✓ *Detecting the Deadlock condition and resolving*
- ✓ *Two-phase locking protocol*
- ✓ *How timestamp-based protocol works*